

During weeks 9–12. This includes Lectures 17–24, and workbooks 9–12.

Meshes

- Lec 17B

- Meshes: collections of triangles

- vertex sharing
- vertex reuse
- index set representation (for faces)

- Good meshes

- consistency of handedness, consistent normals
- Avoid cracking
- avoid T-junctions

- Properties

- Vertex: (barycentric interpolation) vertex colors, vertex normals
- Face: (constant over face) colors, normals

- Splitting

- position is the same
- properties are on faces (different triangles)

- Good triangles:

- not too small
- not too elongated (stretched out)

- Operation/ Representation

- efficient display and storage (index set)
 - compact
 - maps well to hardware: strips/fans, caches, format issues
- Efficient Manipulation

- Fancy data structures:

- keep track of edges
- find neighbors quickly
- polygon soup
- vertex sharing/ indexed representations
- uniform patterns: grids, strips/fans
- Topology (how triangles connect) is separated from vertex info

- Three

- Geometry: basic mesh, list of vertices, faces, simple JS data structures
- BufferGeometry: similar, more efficient, easy transmission to hardware, need to understand buffers first

Update geometry: `geometry.elementsNeedUpdate = true`

Normals

• Lec 17B

Normals: right hand rule, outward facing normal vector, specify if not obvious or face directions

Specify: compute (real), fake: per-face or vertex(interpolate)

Uses: backface culling (do not draw if in the back) and lighting

Lighting

• Lec 17C

Lighting: real-world-like process V.S. Shading: more general term

- local vs. global, the simple models

- Local general: photon arrives from some direction, where to go
 - Absorbed, perfect mirror bounce, micro-geometry, subsurface effect?
 - depends on material and direction
 - Distributions of incoming light and outgoing light (direction and amount)
 - Bi-directional reflectance distribution functions (BRDF): how light is reflected at a surface
 - percent reflected: $r = B(\text{in}, \text{out}, \text{color})$ where in and out are directions
- Local vs global
 - Local: at specific point, each point considered individually, assume light comes from somewhere
 - global: points interact, entire scene considered together, how light transports
- Local lighting: only consider how material reflects light, no global effects(shadows, reflections, spill/bleed)
- Simple local lighting model
 - Assume ideal material
 - make a model
 - tweak the model
 - combine with other models to mix effects

- the components of the basic lighting model (ambient, emission, diffuse, specular)

- Type of models
 - Emission: things give off light
 - Ambient: light from nowhere
 - Specular: direct reflection
 - Diffuse: rough reflection, rough surface

- Diffuse and Specular models and equations

- Diffuse: Ideal materials: lambertian materials p47
 - scatters light in all directions
 - does not matter what direction you look from
 - direction of light matters
 - Reflection:
 - $r_{diffuse} = \hat{n} \cdot \hat{I}$
 - Amount of diffuse reflection = unit surface normal X unit vector to light source
 - $color = (\hat{n} \cdot \hat{I}) c_{light} c_d$ clight: color/intensity of light(RGB), cd: color of the material
- Specular: shiny things, ideal (perfect mirror) p55

- light direction and eye position matter
- eye needs to be in exact correct position
- \hat{e} eye vector and \hat{r} reflection vectors
- Imperfect mirror: phong model
 - graduate fall off as we get away from the optimal direction, keep $\hat{e} \cdot \hat{r} > 0$
 - raise to the power of shininess (phong exponent: p)
 - shinier: more like a perfect mirror: $p \rightarrow \infty$: perfect mirror, p=0: rough
 - $r_{specular} = (\hat{e} \cdot \hat{r})^p$
 - $r_{color} = (\hat{e} \cdot \hat{r})^p c_{light} c_s$ cs: color of material

- Specifying surface (lighting) properties

- phong lighting model
 - $color = c_e + c_a l_a + \sum_{l \in lights} ((\hat{n} \cdot \hat{l}) c_{light} c_d + (\hat{n} \cdot \hat{h})^p c_{light} c_s)$
 - emissive + ambient + diffuse+ specular
- Three.js, phong
 - color
 - color: diffuse and ambient color (main color)
 - specular: specular color
 - Emissive: emissive color
 - Shininess
 - how bright is the shininess (c_s related to c_d)
 - how focused is the shininess (p)
 - Color: plastic: white, metal is colored
 - metalness: metal is more specular, specular keeps color
 - Roughness: dull or shiny, amount of specularity, size of specularity

- When to compute lighting

Use: Once **per pixel** (norm): interpolate normals

Others:

- once per triangle (flat shading)
- once per vertex and interpolate colors (gourand shading)

Texturing

• Lec 18B

- motivations: how to get more than 3 colors on a triangle

Bad ideas

- break into smaller pieces
- colors per pixel

Basic idea: texture mapping

- Define a coordinate system on the triangle, define a map from coordinate to color
- draw: coordinate, look up the color

- steps for “standard” texturing (UV mapping with an image)

Three parts:

- assign coordinates
- loop up the value for the coordinate
- do sth with the value

Basic form: p6

- UV coordinates on triangles (2d)
- Images that define colors (look up color in 2d)
- use looked up color as the color for lighting (for each pixel inside triangle inside graphics hardware)
 - get the barycentric coordinate
 - use this to interpolate the UV values (get UV coordinate)
 - use the UV coordinate to loop up the color
 - use the color as the surface color

- texture use, coordinate ranges

UV in $[0,1]$

outside: p15

- clamp wrapping
- repeat wrapping
- mirror repeat wrapping

Reuse: loading, processing, and storing textures is resource intensive

Steps:

- Specify UVs for triangles
- specify an image to be used for lookup
- specify how to use the image in the material
- Specify all the different parameters (lookp type, wrapping mode.....)

• Lec 18C

– basic textures in THREE

Create objects with UVs p5

- Primitives have predefined UVs
- Geometries,
 - define UVs for each vertex, face, possibility for layers, not part of faces
 - faceVertexUVs: array of layers, each layer is an array of Faces, each face is an array of UVs for each vertex
 - each vertex has a 2d coordinate, each face has 3 vertices, each layer has n faces, each object has 1 layer or more

Load textures p8

- reuse if possible, share between objects
- different kinds of pre-processing happens, not just an image

Attach as colors to objects p9

Others: lightening and material affect color (basicmaterial if no lighting), Y is flipped by default

• Lec 18D

Problems in texture lookups:

- Magnification: a pixel covers less than a pixel in the image
 - basic Sol: nearest neighbor
 - standard sol: Bi-linear interpolation, treat the pixel as a point, does not consider edges of pixels
- Magnification: one pixel covers an area of the texture
 - Basic idea: **Filtering**, average together all the texture pixels the pixel covers
 - problem: what to average and how to average fast

– Filtering

- summed area table, not practice now p16
 - estimate shade as a rectangle
 - pre-compute summed area table
 - fast lookups
 - idea: amortization and pre-computation: save time per-pixel by pre-computing the summed area table
- Mip-maps, modern hardware, fancier p18
 - approximate filter as a square
 - pre-compute multiple sizes of map
 - look up in correct sized maps
 - problem: anisotropy, the areas are not square p20

- anisotropic filtering
 - Hack: use multiple squares
 - Lookup in a smaller image, half by half till 1 pixel lookup
 - loop in smallest image: bi-linear interpolation
 - Lookup in-between images: tri-linear interpolation, require 8 values (4 per layer)
 - historic storage trick
- Summary
 - Pre-compute mip-map (images of various sizes): done with texture loading in practice
 - approximate area with square (Center is easy, size is harder): hardware per pixel
 - figure out which image to sample from: hardware per pixel
 - tri-linear interpolate between levels: hardware per pixel
- Three options p29

• Lec 18E

Step to load color from texture

- make some geometry
- get a picture
- get the picture in the right form
 - square, matches simple geometry, minimal lighting, lots of parts in one image
- assign UV values to vertices
 - Per face: vertex splitting
- enable texturing

Hardware does:

- UV coordinate per pixel
- texture lookup
- texture filtering (Fast: mip maps)

After color: use as the material color for lighting

Visibility

• Lec 19

Process of drawing: triangles, transform (viewing), convert to pixels (rasterization), color the pixel (shading)

Among them: clipping, visibility

viewing: transform from triangles to 2d

rasterization: convert triangles to pixels

Shading (lighting, texturing): color each pixel

- visibility vs. clipping vs. culling

Clipping: decide if the triangle is on the screen

- Skip triangles that are off screen
- clip to the frustum
- includes near and far
- when: happen after the triangle is transformed

Visibility: decide if another triangle blocks the pixel

- having near things block farther ones
- not know until look at all triangles
- immediate mode might draw in any order
- when: depends on the also (often per-pixel after coloring and considering all triangles)
- assume the objects are solid and not-transparent, things are in view
- closer objects occlude farther ones
- Reason: make things look realistic

Culling: skipping primitives based on fast decisions

- When: early
 - discard a whole group of triangles (Triangles in another room)
 - backface culling (After the normal is transformed)
 - ...

- painters algorithm, important in concept, p34

Steps

1. collect all objects
2. sort from back to front
3. draw objects in order (back to front)

Problems

- need all objects to sort, not immediate mode
- what happens with ties
 - cutting objects
- what happens with intersecting objects
- Inefficiency: resort when camera moves
 - use good data structure (BST)
- Inefficiency: draw things that get covered

- Z-buffer, used in practice

Goal: order independent, **immediate mode**(1 triangle at a time)

Idea: store the depth per-pixel

Use: an extra number per-pixel, color buffer(RGB), z-buffer(Z)

General

- start with all pixels at max distance
- anything that will be drawn will be in front of the background
- why the far distance of the camera is important
- Historically: limited precision, 16 bits, avoid making far too big

Steps

1. draw pixel x,y with color c and depth z
2. frame buffer FB, Z-Buffer (ZB)
3. Old- no Z: white pixel: $FB(x,y) = c$
4. New - Z (read/test/write)
 1. $pz = ZB(x,y)$
 2. if $z < pz$:
 1. $FB(x,y) = c$
 2. $ZB(x,y) = z$

Order independent

Problem

- Requires memory
- require read/modify/write
 - memory reads can be slow
 - need to wait until prior writes are done
- requires storing distances, (old: 16 bits, scaling issue)
- Tie
 - Decide if $<$ or \leq , order matters
- **Z-fighting**: two objects are very similar distances
 - might be a tie (drawing order matters)
 - numerical noise might break the tie
 - flipping between close objects
- Triangle partially fills a pixel
 - z-buffer is yes/no
 - does not consider partial iflling
 - Aliasing:
 - anti-alising the edges fo the objects is hard
- Overdraw efficiency

- throw away pixel after it is computed
- after we have computed the color
- Wasted effort
- overwrite is even more wasteful than Z-fail
- Future: ways to avoid
- Semi-transparent objects
 - back object need to be drawn first
 - need to be blend with correct background
 - close object prevents objects behind from being drawn
 - Solution:
 - sort objects (painters algo)
 - transparent objects first

Summary: clear z-buffer to max distance, replace pixel write with read/test/write

Good:

- simple,
- Use memory now cheap
- generally order independent
- do not need objects ahead of time
- easy to implement in hardware

Problems

- sometime order matters
- can have efficiency issues
- z-fighting problems
- aliasing
- do not handle transparency

● Lec 20

Interactive rendering:

- every triangle is independent
- every triangle gets lit, shaded, textured, visibility checked...
- fast, all effects possible

Graphics Pipeline, hardware

- Steps

1. model objects (make triangles)
2. transform (find point and positions)
3. shade (lighting – per triangle/vertex)
4. Transform (projection)
5. rasterize (figure out pixels)
6. Shade (per pixel coloring)
7. write pixels (with z buffer)

Drawing process:

- Modeling (model transformation), shading, viewing, clipping ,projecting, visibility, rasterization

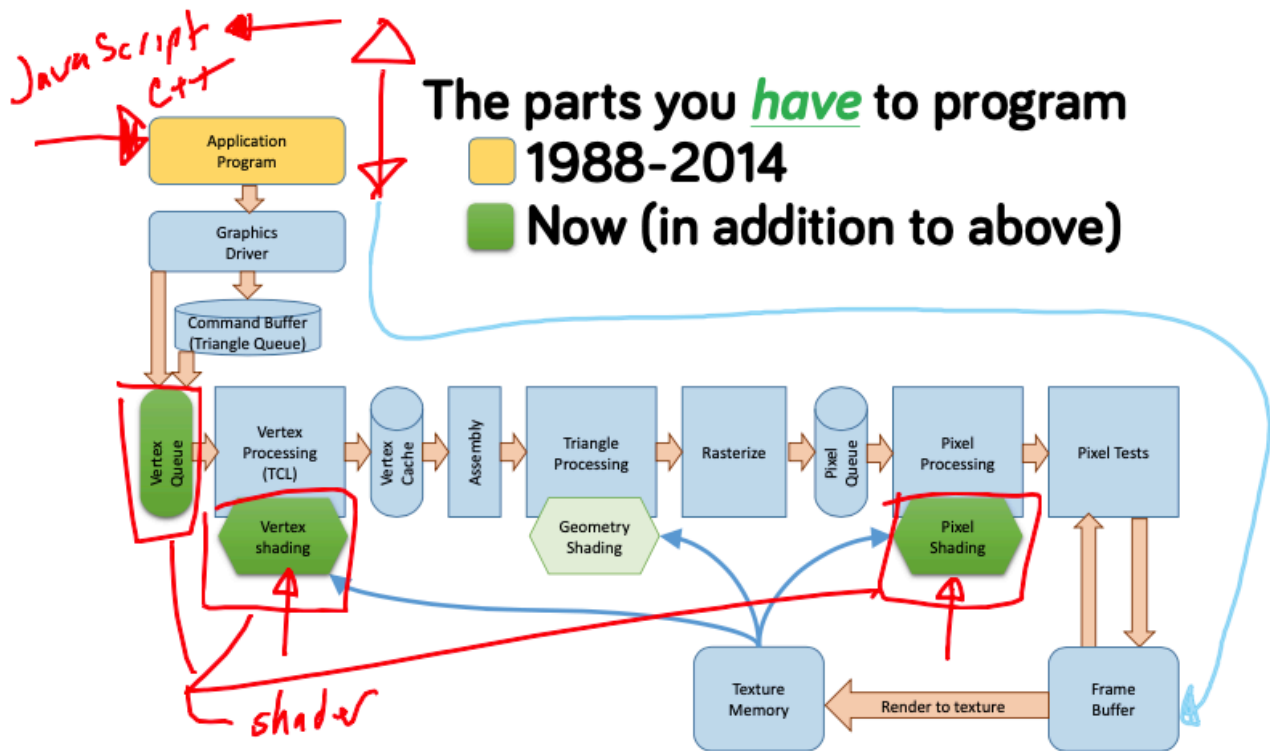
In the pipeline

- Triangles are independent
 - no pipe stall (the param in the next step is defined in the current step p8)
 - no complexity of handling stalls
- vertices are independent
 - parallelize, process multiple triangles, pixels at the same time
- pixels in the triangles are independent
- do not need to finish 1 before start 2

Parallelization: fast, require specific model in hardware, program interface is designed for the model

- vertex operations
 - split triangles/re-assemble
 - compute per-vertex not per triangle
- pixel(fragment) operations
 - lots of potential parallelism
 - less predictable
- use queues and caches

History: p33



Steps details:

- draw the triangle
 - per vertex, not per triangle
 - allows sharing vertices between triangles, or make all same to get truly flat
 - position, color, normal in vertices
- command buffer
 - split triangles into vertices
- vertex queue: buffer/queue for the vertices to process vertices in parallel
- vertex processing (TCL)
 - Compute transform – compute x' and n' on the screen, clip (old days), Light–compute c' per vertex
 - Add info to vertices
 - independent vertex
 - for each vertex
 - in: vertex info x, n, c + global information (current transform, lighting)
 - out: vertex info: x, n, c, x', n', c' (vertex lit color)
- Vertex cache:
 - store several processed vertices in the cache so that we can re–use if triangles share the vertex
 - Before is important to performance, now not
- Assembly: put triangles back together, triangles exist
- Triangle processing: in the fixed–function pipeline, not this step (maybe clipping)
- Rasterizer: convert triangles to a list of pixels
 - In: triangle with values ($x', c'...$)

- pixels with values (x' , c')
- figure out which pixels a primitive covers
- turns primitives into pixels
- interpolate vertex values to interior pixels, barycentric coordinates
 - pixel values from 3 vertices, all are interpolated
- Pixel queue: to process pixels in parallel
- pixel processing: usually color, sometimes z, not in 1988 pipeline
 - pixel in, pixel out, each independent, just change values (re-compute)
 - not change position (make another pixel)
 - change or reject other values
- Pixel tests
 - z-buffer, alpha/color/stencil-tests, color blending
- Read/write cycle from pixel tests to frame buffer
 - needed for each pixel
 - potential memory bandwidth **bottleneck** (problem)
 - failed z-test: removed by z-buffer
- Frame buffer: sent to the screen

Pixel and fragment

- Pixel: a point on the screen
- fragment: a dot on the triangle
 - might not become a pixel (fails z-test), might only be part of a triangle

Summary:

- transfer info about vertices from application
- process vertices
- assemble/rasterize triangles
- process fragments/pixels
- Test/write pixel results

- how does this lead to the shader model and efficiency issues

Shader Basics

• Lec 21

program do:

- set up for drawing
- repeat for every frame (Render)
 - clear the screen
 - for each object: draw a group of triangles

Object: a group of triangles sent together

- the vertices
- how they connect to triangles
- what shaders to use in processing them
- what parameters to use across all of them

limited in what we can change within a group

Constant state: can not change certain things while drawing a triangle, extends to group of triangles

Uniform over group of triangles: camera, frame buffer, light, texture maps....

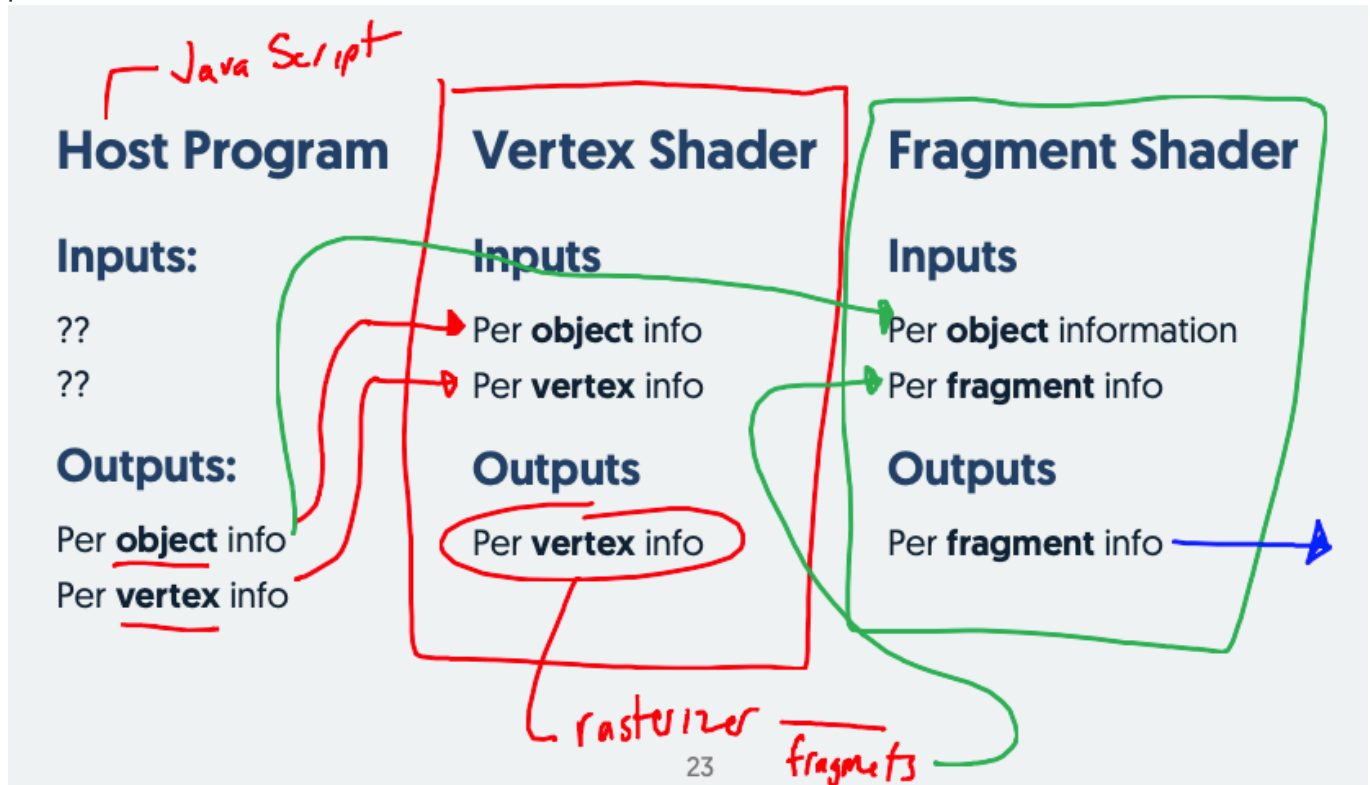
Three.js: scene and material with common properties

Can specify:

- global info: whole screen, window
- per triangle group info: uniforms
- per vertex info: position, color, normal...
- **NOT** per pixel

– Programmable shading into the graphic pipeline

Two main parts



vertex processor

- Input: vertex with info (attributes from the host program), output vertex with more info (more attributes)
- before: position, normal, color, UV
- After: position, normal, color, UV, computed: screen space position, vertex lit color, screen space normal
- Needed: screen space position for rasterizer/z-test (with other we want for coloring)
- For rasterizer:
 - use the screen coordinate to generate list of fragments
 - interpolate other values so each fragment has the properties

fragment processor

- have (depend on three vertices)
 - screen space position
 - interpolated values (attributes)
- we do:
 - each fragment is processed independently with info from rasterizer
 - figure color and depth
- Before: screen space position, depth (z) value, other interpolated properties
- After: **same** screen space position, depth value, color to write to frame buffer

- Shading language GLSL

Shaders: write in shading language

- vertex shaders: compute new attributes for vertices, `gl_position`
- Fragment shaders: compute color/depth for fragments, `gl_fragcolor`

History, p26

GLSL: runtime compilation

- parts: shading language, host api (JS)
- C-like, C syntax, strict typing, operator overloading
- math data types (metrics, vectors), built-in functions
- features for tying things together
- each shader has own program
- connect to other parts via variables, like global variables, special declarations, few built ins

- variable types (how shaders communicate)

Uniform

- host triangle groups (JS)
- not changed inside the triangle group
- some built in
- JS: creating mapping of names, convert types, provide mapping table (Three do this)
- P52

Uniforms

```
{
  uniforms: {
    var1 : { value: 10.0},
    var2 : { value: new T.Vec2(1,2);}
  }
}
```

uniforms are one parameter to `ShaderMaterial`
provide a dictionary that maps **variables** to **values**
each value is a dictionary (with the key **value**)

Built-In Uniforms

Some of the uniforms are "built in"

```
uniform mat4 modelMatrix;
uniform mat4 modelViewMatrix;
uniform mat4 projectionMatrix;
uniform mat4 viewMatrix;
uniform mat3 normalMatrix;
uniform vec3 cameraPosition;
```

These are added to your vertex program

Some of these are added to your fragment program

Attribute p54

- Per vertex input, host program
- Position, color, normal, UV
- transfer as block of memory (buffers)
- Three: position, normal, uv

Using Attributes and Uniforms

From the THREE documentation

<https://threejs.org/docs/#api/en/renderers/webgl/WebGLProgram>

```
gl_Position = projectionMatrix * modelViewMatrix * vec4( position, 1.0 );
```

or alternatively

```
gl_Position = projectionMatrix * viewMatrix * modelMatrix * vec4( position, 1.0 );
```

Varying

- per vertex output from vertex shader and input to fragment shader
- interpolate from 3 vertices
- gl_position

Three: ShaderMaterials

– connection to the host program

Just strings, separate files and loaded

Steps: compile, link, use

• Lec 22

– lighting in shaders

Where

- vertex shader
 - compute lighting at every vertex, interpolate colors across triangle
 - per-vertex lighting, gouraud shading
 - old days
- fragment shader
 - interpolate normals and surface colors
 - compute lighting at every pixel
 - per-pixel shading
 - hardware is fast

Same function as before (lec 17), workbook 10

codes: p40

Shader Programming

- shader idioms (using built-ins)

more efficient and convenient

Mix(a, t1, t2): $a * t1 + (1-a) * t2$

step, fact, sign, min, max

- shader anti-aliasing, Lec 21 p65, Lec 22 p12

smoothstep: one **sided** only for 0-1 transition

Df dx: derivative of function wrt x

Fwidth: derivative of function wrt x and y

smoothstep handles texture magnification

no texture minification

- need filtering, hard to do in a general way for procedures: over sampling

Easier to do for images

• Lec 22

Noises, randomness p24

- pseudo-random
 - pattern too complex to see
 - still controlled/ deterministic
- structured
 - control properties that we care about
- Simple method: $color = f(u), r = \text{fract}(\sin(u) * 10000)$
 - problem: aliasing
 - use interpolate to make smooth
- Perlin noise
 - classic noise function, with more efficient GPU
 - Noise-controlled pseudo-randomness
 - coherence at different frequencies

- displacement map shading, p44

Fake shape with vertex shader, displace points

Use dots to move things in the normal position

- only move vertices, plane the same
- did not change lighting
- spacing depends on mapping
- some triangles get really stretched
- do change the shape

Bad:

- depends on meshing, need small triangles
- messes up the mesh
- need to get things aligned with vertex samples
- may want to make more triangles

Texturing Tricks

- fake normals for fake lighting, p49

Flat

- make many triangles
- fake it with texture: smooth with lighting

Not in the right shape, just make it looks right with shape

Lec 24

- bump and normal maps

p11

Compare to displacement maps

- no change to geometry, no shadows, occlusions, vertex or fragment
- both just change normals but in different ways

Normal maps

- texture lookup of normal vector, RGB \rightarrow xyz directions, middle value $(128/256) = 0$
- need to renormalize
- vector relative to real normal
- deal with xyz of normals

Bump maps

- deal with the height of changes in the normal

Good

- easy to specify surface details
- does not change shapes
- basic lighting effects
- Works iwth lighting
- easy in three

Bad

- not change side view
- not cause occlusions
- not for big effects
- No shadows
- need reflection

- skyboxes, p21

Far away: position not matter, orientations do

a box for background,

always styas centered at the camerra, move with the camera

help reflections:

- assume objects being looked at are far away
- view direction matters
- point position doesn't

- environment mapping

cube or sphere (around object)

position does not matter, normal matters

assume point is at center of sphere/cube

Lookup direction is based on viewing direction

Require:

- small object/far environment
- eye is far (direction amtter)
- position do not matter
- small curvy objects good, large flat bad

- spherical and cubic representations for maps

sphere

- sampling issues at pole,
- single images
- capture in 1 photograph

Cube

- sampling issues in corners
- Images are human viewable
- maps nicely to graphics hardware

- layered textures

The environment map is light, real scene to create brightness, not limited to small set of light sources

Often states (pre-compute ahead of time)

Dynamic environment map p34 multi-pass rendering

- draw the scene once or multiple times, remember info
- draw the scene again with the info
- both steps when world changes, step 1 first
- store info in textures
- no need to draw everything in step1

- light maps

Scalling problem p40

- Book matching, flipped Y, repeat wrapping

Light maps: put lighting into the texture p48

Good:

- can be fancy
- pre-computed

Bad

- must be known ahead
- can not change when camera moves, object moves

Special case

- self shadowing: importing for conveying shape, pretend light comes from all directions (ambient)
- amount each point is visible
- corners and crevices are dark

Ambient Occlusion shading (all directions for objects, blocks, light)

- pre-compute for all points with special tools or clever hacks
- like a light-map, ambient occlusion map

- shadow maps p52

non-local

only one way to do shadows

Idea: can light see the object? -> put "camera" from light source

Step

1. take a picture from light position
 - camera at light source
 - what objects are visible to the light (others are shadow)
 - use this picture as the shadow map
2. draw the regular picture
 - for each pixel on an object
 - see if pixel is visible in shadow map

Test: check the pixel is the same pixel in the shadow map

- check color (no)
- Common: check depth
 - Problem: small errors will lead to larger
 - shadow map is a z-buffer (depth map)

Resolution

size matters

- Spotlight (Can be small)
- directional light source (Area)
- point light? (needs to be a cube/sphere)